

Parallel computation of the sequence of iterates of a function

Jean-Michel Fourneau Maël Guiraud Yann Strozecki

December 22, 2017

Abstract

1 Model

Let \mathcal{S} be a finite set of states and let f be a computable function from \mathcal{S} to \mathcal{S} . The aim of this short article is to propose practical methods to compute the sequence $(s, f(s), \dots, f^n(s))$ in parallel. When it is possible to compute efficiently $f^i(s)$ using only s and i , it is easy to compute the sequence in parallel by assigning one $f^i(s)$ to each processor. However, for general f we need $f^i(s)$ to compute $f^{i+1}(s)$ and there is no simple scheme to parallelize the computation of the sequence. **TODO: Is there a link with PPAD ?** Hence, we consider additional structure on \mathcal{S} to make the problem tractable. The set \mathcal{S} is equipped with a partial order \leq and it has a smallest element \perp and a greatest element \top . We assume that f is monotone, that is if s_1 and s_2 are elements of \mathcal{S} such that $s_1 \leq s_2$ then $f(s_1) \leq f(s_2)$.

TODO: Define a simple model of parallelization + give a name to the problem

This models a concrete problem: the simulation of a random process. In that setting, the function f which describe the dynamic of the process also depends on the value of some random variable. In a simulation we use a random seed and then a random generator to generate the sequence of values of the random variable from the random seed. Say that the seed is a m bit integers in $[2^m]$, we have a pseudo random generator R which maps $[2^m]$ to $[2^m]$. A state of the system is $(s, y) \in \mathcal{S} \times [2^m]$ where s describes the state of the random process and y is the current random number. A transition of the random process is given by a function f from $\mathcal{S} \times [2^m]$ to \mathcal{S} thus a pair (s, x) is mapped to $(f(s, x), R(x))$ and we want to compute the iterates by this function.

2 Method

The main idea of the method is to divide $[n+1]$ into intervals of size t , on which the sequence of iterates will be computed independently. We denote by I_j the interval $\{tj, \dots, tj+t-1\}$. Each processor will be assigned the task to compute the sequence on some interval I_j that is $(f^i(s))_{i \in I_j}$. We assume that we have

some central memory in which the final solution is stored and to which each processor can write.

2.1 Two bounds

We describe an algorithm which produces the sequence $(f^i(s))_{i \in [n+1]}$. Each interval I_j has one of three states during the algorithm: ToCompute, InProgress and Done and we store for each interval I_j the states s_j^{min} and s_j^{max} . The algorithm is the following: at the beginning, all intervals are marked ToCompute, for all $j > 0$, $s_j^{min} = \perp$, $s_j^{max} = \top$ and $s_0^{min} = s_0^{max} = s$. In parallel, select a free processor P and an interval I_j in state ToCompute. The processor P first and set the state of I_j to InProgress and then computes the two sequences $(f^i(s_j^{min}))_{i \in [t]}$ and $(f^i(s_j^{max}))_{i \in [t]}$ iteratively. When the sequences are computed, we have access to $f^t(s_j^{min})$ and $f^t(s_j^{max})$. If $f^t(s_j^{min}) > s_{j+1}^{min}$ or $f^t(s_j^{max}) < s_{j+1}^{max}$ then better bounds have been found and P sets $s_{j+1}^{min} = f^t(s_j^{min})$, $s_{j+1}^{max} = f^t(s_j^{max})$ and the state of I_{j+1} to ToCompute. Finally, if s_j^{min} is equal to s_j^{max} , the result of the simulation is stored as the solution on the interval I_j and I_j is set to Done.

TODO: écrire l'algo dans un environnement algo

Remark that the choice of an interval with state ToCompute by a free processor is not specified. In practice we propose two ways to select it. The first is to chose the interval of smallest index. The second is to cut $[n+1]$ into l consecutive intervals (containing several I_j), where l is the number of processors. Then we affect each processor to an interval and when it is free, we affect the smallest I_j with state ToCompute in this interval. **TODO: Etre bien plus précis dans les politiques de choix d'intervalle surtout pour le deuxième algo et lui donner un nom**

TODO: Faire un petit modèle probabiliste qui pour une proba de coupler donnée combien on va faire de calculs. Ça serait bien de se servir de ça pour couper des intervalles de la bonne taille.

Theorem 1. *Algorithm TwoBounds produces the solution of the problem blah in a time bounded by the one of the obvious sequential algorithm.*

Proof. We first prove that the algorithm terminates. At any point in time, there is a processor which is computing a part of the final solution. By induction on the time, assume a processor P finishes to compute the sequence on an interval. Then consider the smallest interval I_j which is not marked Done. If it is in state InProgress, then a processor is computing the solution on I_j . If it is in state ToCompute, it will be selected by the free processor P and since I_{j-1} is in state Done, $x_j^{min} = x_j^{max}$ and P will compute the solution on I_j .

We must also prove that when an interval is set to Done, the right sequences of states has been computed. It follows from the fact that $x_j^{min} \leq f^{jt}(s)$ and $x_j^{max} \geq f^{jt}(s)$. It can be proved by induction on the number of times the variables x_j^{min} and x_j^{max} are updated, using the monotonicity of f and the initialization of these variables to \top and \perp . Hence when $x_j^{min} = x_j^{max}$ then it is also equal to $f^{jt}(s)$ and this value is enough to compute the sequence on I_j . \square

TODO: Donner une borne simple en moyenne quand les intervalles sont suffisamment long

Remark that all processors do the simulation using only their private memory and the two states at the beginning of the interval. This scheme is thus adapted to a distributed computing environment where the cost of transmitting information between processors is high.

In our application to the simulation of a random process, the states of the process are often equipped with a partial order, but the random integers are not. However, we have the following property in our random system, if $s_1 \leq s_2$ then $f(s_1, x) \leq f(s_2, x)$. In fact, the random value alone is used to select an action to apply to the system and the action does not depends on the state of the system.

Since a random integer in the sequence depends only on the previous one, we can do the following trick to transform our problem. Instead of using a single seed to generate all pseudo random values, use one seed for each interval I_j . That is we have a collection of seeds $x_0, \dots, x_{n/t}$ and the i th random number will be equal to $R^{i \bmod t}(x_{i/t})$ instead of $R^i(x)$. By doing that, we change the sequence we simulate, but it is still the same random process, that is all realizations of the random process still occur with the same probability.

TODO: Mieux expliquer cette partie

2.2 Algos avec intervalles de taille variables

Inspiré de l’algo précédent mais plus souple. On coupe le temps en autant d’intervalles que de processeurs. Ils sont simulés indépendamment et on note quand ça couple et on renvoie la partie bien simulée, les seeds de la parties non couplée (seulement une petite partie). On réaffecte les processeurs sur les parties non encore trouvées

2.3 One bound

3 Experiments

The random process we use in our experiments is the following one. We have a system composed of m finite queues of capacity `BUFF_MAX` in tandem. A queue is characterized by its number of client C_i . Every queue i has tree events that can occurs:

- Arrival, C_i is increased by one, if it is not already to `BUFF_MAX`.
- Service, a client leave the queue i and goes in the queue $i + 1$. C_i is decreased by one and C_{i+1} is increased by one. C_{i+1} is equal to `BUFF_MAX`, the client is lost.
- Departure, the client leave the system. The number of client of the queue is decreased by one.

For the queue i , every event have a probability denoted by respectively a_i , s_i and d_i . The last queue has no service, thus $s_m = 0$. There is thus a total of $3m - 1$ different events that can change the system.

The sequence that we want to parallelize is a succession of drawing of one of those $3m - 1$ events. Every event is drawn randomly, following a pseudo random generator.




Figure 1: A system with 3 queues

TODO: Precise description of the random process avec un dessin

Our experimentations are made with the following set-up. The program master is running on a MacBook Air, with a processor 2.2 GHz Intel Core i7, and 8 Go of RAM DDR3 at 1600 MHz. The operating system of the machine is macOS High Sierra v10.13.1. The source code is compiled with gcc version 7.1.0 (Homebrew GCC 7.1.0 –without-multilib).

We use up to 7 servers are running on Raspberry Pi 3 , Model B, with 1GB of RAM. Their operating system is Raspbian GNU/Linux 8.0, installed on a micro SD card element14 with a size of 8GB. The source code is compiled with gcc version 4.9.2 (Raspbian 4.9.2 – 10).

All the machines are connected on by a local network through an HP 14-10 8G switch.

TODO: Description of the experimental settings: quel système d’exploitation/matériel/compilateur/librairie

TODO: Experiments avec des jolies courbes et des interprétations

Practical problems: cost of the network transmission, especially for transmitting long sequences. -¿ measure the time of a two way trip for a small message and the time of sending an interval. To say that in practice we will not compute the whole sequence but statistic on it which could help reduce the use of the network.